

Data Parallelism

Different Data Slices





Data 1

Data 2


Data 3

Data 4

Same Model
Architecture



Aggregate and Update model
variables



New
Master
Model

Data Parallelism

Different Data Slices

Data 1

Data 2






Data 3

Data 4

Same Model Architecture

Aggregate and Update model variables

New Master Model



Data Parallelism

Different Data Slices





Data 1

Data 2


Data 3

Data 4

Same Model Architecture



Aggregate and Update model variables



New Master Model

`tf.distribute.Strategy`

- High-level APIs
- Custom training loops
- TensorFlow 2: eager mode & graph mode
- Supported on multiple configurations.
- Convenient to use with little to no code changes

Commonly used terms

- Device
- Replica
- Worker
- Mirrored variable

Commonly used terms

- Device



CPU

Accelerator: GPU, TPU

- Replica
- Worker
- Mirrored variable

Commonly used terms


- Device



CPU

Accelerator: GPU, TPU

- Replica



- Worker

- Mirrored variable

Commonly used terms


- Device



CPU

Accelerator: GPU, TPU

- Replica



- Worker



- Mirrored variable

Commonly used terms



- Device
- Replica
- Worker

- Mirrored variable



CPU

Accelerator: GPU, TPU



Classifying strategies

Hardware platforms

- Single-machine multi-device
- Multi-machine (with 0 or more accelerators)

Training

- Synchronous (All-reduce)
- Asynchronous (Parameter Server)

Classifying strategies

Hardware platforms

- Single-machine multi-device
- Multi-machine (with 0 or more accelerators)




Training

- Synchronous (All-reduce)
- Asynchronous (Parameter Server)

Classifying strategies

Hardware platforms

- Single-machine multi-device
- Multi-machine (with 0 or more accelerators)



Training

- Synchronous (All-reduce)
- Asynchronous (Parameter Server)

Classifying strategies

Hardware platforms

- Single-machine multi-device
- Multi-machine (with 0 or more accelerators)



Training

- Synchronous (All-reduce)
- Asynchronous (Parameter Server)

Classifying strategies

Hardware platforms

- Single-machine multi-device
- Multi-machine (with 0 or more accelerators)



Training

- Synchronous (All-reduce)
- Asynchronous (Parameter Server)

MirroredStrategy

MultiWorkerMirroredStrategy

ParameterServerStrategy

DefaultStrategy

TPUStrategy

CentralStorageStrategy

OneDeviceStrategy

MirroredStrategy

- Single-machine multi-GPU
- Creates a replica per **GPU**
- Each variable is ***mirrored***
- All-reduce ***across devices***

MultiWorkerMirroredStrategy

ParameterServerStrategy

DefaultStrategy

TPUStrategy

CentralStorageStrategy

OneDeviceStrategy

MirroredStrategy

- Single-machine multi-GPU
- Creates a replica per **GPU**
- Each variable is ***mirrored***
- All-reduce ***across devices***

MultiWorkerMirroredStrategy

ParameterServerStrategy

DefaultStrategy

TPUStrategy

- Same as **MirroredStrategy**
- All-reduce across ***TPU cores***

CentralStorageStrategy

OneDeviceStrategy

MirroredStrategy

- Single-machine multi-GPU
- Creates a replica per **GPU**
- Each variable is **mirrored**
- All-reduce **across devices**

MultiWorkerMirroredStrategy

- Multi-machine multi-GPU
- Replicates variables per device **across workers**
- All-reduce based on
 - hardware
 - network topology
 - tensor sizes

ParameterServerStrategy

DefaultStrategy

TPUStrategy

- Same as **MirroredStrategy**
- All-reduce across **TPU cores**

CentralStorageStrategy

OneDeviceStrategy

MirroredStrategy

- Single-machine multi-GPU
- Creates a replica per **GPU**
- Each variable is **mirrored**
- All-reduce **across devices**

MultiWorkerMirroredStrategy

- Multi-machine multi-GPU
- Replicates variables per device **across workers**
- All-reduce based on
 - hardware
 - network topology
 - tensor sizes

ParameterServerStrategy

DefaultStrategy

TPUStrategy

- Same as **MirroredStrategy**
- All-reduce across **TPU cores**

CentralStorageStrategy

- Variables are **not mirrored**
(instead placed on the CPU)
- Done in-memory on a device

OneDeviceStrategy

MirroredStrategy

- Single-machine multi-GPU
- Creates a replica per **GPU**
- Each variable is **mirrored**
- All-reduce **across devices**

MultiWorkerMirroredStrategy

- Multi-machine multi-GPU
- Replicates variables per device **across workers**
- All-reduce based on
 - hardware
 - network topology
 - tensor sizes

ParameterServerStrategy

- Some machines designated as **workers**
- Some others as **parameter servers**

DefaultStrategy

TPUStrategy

- Same as **MirroredStrategy**
- All-reduce across **TPU cores**

CentralStorageStrategy

- Variables are **not mirrored**
(instead placed on the CPU)
- Done in-memory on a device

OneDeviceStrategy

MirroredStrategy

- Single-machine multi-GPU
- Creates a replica per **GPU**
- Each variable is **mirrored**
- All-reduce **across devices**

MultiWorkerMirroredStrategy

- Multi-machine multi-GPU
- Replicates variables per device **across workers**
- All-reduce based on
 - hardware
 - network topology
 - tensor sizes

ParameterServerStrategy

- Some machines designated as **workers**
- Some others as **parameter servers**

TPUStrategy

- Same as **MirroredStrategy**
- All-reduce across **TPU cores**

CentralStorageStrategy

- Variables are **not mirrored**
(instead placed on the CPU)
- Done in-memory on a device


DefaultStrategy

- Simple Passthrough

OneDeviceStrategy

- Single device

Mirrored Strategy



Mirrored Strategy

- Model declaration
- Data preprocessing

```
model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(10)
])
```

```
model.compile(
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    optimizer=tf.keras.optimizers.Adam(),
    metrics=['accuracy'])
```



```
model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(10)
])
```

```
model.compile(
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    optimizer=tf.keras.optimizers.Adam(),
    metrics=['accuracy'])
```

```
model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(10)
])
```

```
model.compile(
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    optimizer=tf.keras.optimizers.Adam(),
    metrics=['accuracy'])
```

```
def scale(image, label):  
    image = tf.cast(image, tf.float32)  
    image /= 255  
    return image, label
```

```
num_train_examples = info.splits['train'].num_examples  
num_test_examples = info.splits['test'].num_examples
```

```
BUFFER_SIZE = 10000
```

```
BATCH_SIZE = 64
```

```
train_dataset = mnist_train.map(scale).cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)  
eval_dataset = mnist_test.map(scale).batch(BATCH_SIZE)
```

```
def scale(image, label):  
    image = tf.cast(image, tf.float32)  
    image /= 255  
    return image, label
```

```
num_train_examples = info.split('train').num_examples  
num_test_examples = info.split('test').num_examples
```

```
BUFFER_SIZE = 10000
```

```
BATCH_SIZE = 64
```

```
train_dataset = mnist_train.map(scale).cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)  
eval_dataset = mnist_test.map(scale).batch(BATCH_SIZE)
```

```
def scale(image, label):  
    image = tf.cast(image, tf.float32)  
    image /= 255  
    return image, label
```

```
num_train_examples = info.splits['train'].num_examples  
num_test_examples = info.splits['test'].num_examples
```

```
BUFFER_SIZE = 10000
```

```
BATCH_SIZE = 64
```

```
train_dataset = mnist_train.map(scale).cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)  
eval_dataset = mnist_test.map(scale).batch(BATCH_SIZE)
```

```
strategy = tf.distribute.MirroredStrategy()
```

```
print('Number of devices: {}'.format(strategy.num_replicas_in_sync))
```

```
strategy = tf.distribute.MirroredStrategy()
```

```
print('Number of devices: {}'.format(strategy.num_replicas_in_sync))
```

```
def scale(image, label):  
    image = tf.cast(image, tf.float32)  
    image /= 255  
  
    return image, label  
  
num_train_examples = info.splits['train'].num_examples  
num_test_examples = info.splits['test'].num_examples  
  
BUFFER_SIZE = 10000  
  
BATCH_SIZE_PER_REPLICA = 64  
BATCH_SIZE = BATCH_SIZE_PER_REPLICA * strategy.num_replicas_in_sync  
  
train_dataset =  
mnist_train.map(scale).cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)  
eval_dataset = mnist_test.map(scale).batch(BATCH_SIZE)
```



```
def scale(image, label):  
    image = tf.cast(image, tf.float32)  
    image /= 255  
  
    return image, label  
  
num_train_examples = info.splits['train'].num_examples  
num_test_examples = info.splits['test'].num_examples  
  
BUFFER_SIZE = 10000  
  
BATCH_SIZE_PER_REPLICA = 64  
BATCH_SIZE = BATCH_SIZE_PER_REPLICA * strategy.num_replicas_in_sync  
  
train_dataset =  
mnist_train.map(scale).cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)  
eval_dataset = mnist_test.map(scale).batch(BATCH_SIZE)
```

```
def scale(image, label):  
    image = tf.cast(image, tf.float32)  
    image /= 255  
  
    return image, label  
  
num_train_examples = info.splits['train'].num_examples  
num_test_examples = info.splits['test'].num_examples  
  
BUFFER_SIZE = 10000  
  
BATCH_SIZE_PER_REPLICA = 64  
BATCH_SIZE = BATCH_SIZE_PER_REPLICA * strategy.num_replicas_in_sync  
  
train_dataset =  
mnist_train.map(scale).cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)  
eval_dataset = mnist_test.map(scale).batch(BATCH_SIZE)
```

```
with strategy.scope():  
    model = tf.keras.Sequential([  
        tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),  
        tf.keras.layers.MaxPooling2D(),  
        tf.keras.layers.Flatten(),  
        tf.keras.layers.Dense(64, activation='relu'),  
        tf.keras.layers.Dense(10)  
    ])  
  
model.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
              optimizer=tf.keras.optimizers.Adam(),  
              metrics=['accuracy'])
```

```
with strategy.scope():
```

```
    model = tf.keras.Sequential([  
        tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),  
        tf.keras.layers.MaxPooling2D(),  
        tf.keras.layers.Flatten(),  
        tf.keras.layers.Dense(64, activation='relu'),  
        tf.keras.layers.Dense(10)  
    ])
```

```
model.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
              optimizer=tf.keras.optimizers.Adam(),  
              metrics=['accuracy'])
```

```
with strategy.scope():  
    model = tf.keras.Sequential([  
        tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),  
        tf.keras.layers.MaxPooling2D(),  
        tf.keras.layers.Flatten(),  
        tf.keras.layers.Dense(64, activation='relu'),  
        tf.keras.layers.Dense(10)  
    ])
```

```
model.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
              optimizer=tf.keras.optimizers.Adam(),  
              metrics=['accuracy'])
```

Epoch 1/12

INFO:tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadcast to
INFO:tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadcast to
INFO:tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadcast to
INFO:tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadcast to

Training across local GPUs

`tf.distribute.MirroredStrategy`

- Each variable in the model is mirrored across all replicas
- Variables are treated as `MirroredVariable`
- **Synchronization** done with NVIDIA NCCL

```
# Create Datasets from the batches
```

```
train_dataset = tf.data.Dataset.from_tensor_slices((train_images, train_labels))  
                .shuffle(BUFFER_SIZE).batch(GLOBAL_BATCH_SIZE)
```

```
test_dataset = tf.data.Dataset.from_tensor_slices((test_images, test_labels))  
                .batch(GLOBAL_BATCH_SIZE)
```

```
# Create Distributed Datasets from the datasets
```

```
train_dist_dataset = strategy.experimental_distribute_dataset(train_dataset)  
test_dist_dataset = strategy.experimental_distribute_dataset(test_dataset)
```



```
# Create Datasets from the batches
```

```
train_dataset = tf.data.Dataset.from_tensor_slices((train_images, train_labels))  
                .shuffle(BUFFER_SIZE).batch(GLOBAL_BATCH_SIZE)
```

```
test_dataset = tf.data.Dataset.from_tensor_slices((test_images, test_labels))  
                .batch(GLOBAL_BATCH_SIZE)
```

```
# Create Distributed Datasets from the datasets
```

```
train_dist_dataset = strategy.experimental_distribute_dataset(train_dataset)
```

```
test_dist_dataset = strategy.experimental_distribute_dataset(test_dataset)
```

```
# Create Datasets from the batches
```

```
train_dataset = tf.data.Dataset.from_tensor_slices((train_images, train_labels))  
                .shuffle(BUFFER_SIZE).batch(GLOBAL_BATCH_SIZE)
```

```
test_dataset = tf.data.Dataset.from_tensor_slices((test_images, test_labels))  
                .batch(GLOBAL_BATCH_SIZE)
```

```
# Create Distributed Datasets from the datasets
```

```
train_dist_dataset = strategy.experimental_distribute_dataset(train_dataset)  
test_dist_dataset = strategy.experimental_distribute_dataset(test_dataset)
```

```
# Create Datasets from the batches
```

```
train_dataset = tf.data.Dataset.from_tensor_slices((train_images, train_labels))  
                .shuffle(BUFFER_SIZE).batch(GLOBAL_BATCH_SIZE)
```

```
test_dataset = tf.data.Dataset.from_tensor_slices((test_images, test_labels))  
                .batch(GLOBAL_BATCH_SIZE)
```

```
# Create Distributed Datasets from the datasets
```

```
train_dist_dataset = strategy.experimental_distribute_dataset(train_dataset)  
test_dist_dataset = strategy.experimental_distribute_dataset(test_dataset)
```

```
EPOCHS = 10
for epoch in range(EPOCHS):
    # Do Training
    total_loss = 0.0
    num_batches = 0
    for batch in train_dist_dataset:
        total_loss += distributed_train_step(batch)
        num_batches += 1
    train_loss = total_loss / num_batches
```

```
EPOCHS = 10
for epoch in range(EPOCHS):
    # Do Training
    total_loss = 0.0
    num_batches = 0
    for batch in train_dist_dataset:
        total_loss += distributed_train_step(batch)
        num_batches += 1
    train_loss = total_loss / num_batches
```

```
EPOCHS = 10
for epoch in range(EPOCHS):
    # Do Training
    total_loss = 0.0
    num_batches = 0
    for batch in train_dist_dataset:
        total_loss += distributed_train_step(batch)
        num_batches += 1
    train_loss = total_loss / num_batches
```

```
EPOCHS = 10
for epoch in range(EPOCHS):
    # Do Training
    total_loss = 0.0
    num_batches = 0
    for batch in train_dist_dataset:
        total_loss += distributed_train_step(batch)
        num_batches += 1
    train_loss = total_loss / num_batches
```

```
EPOCHS = 10
for epoch in range(EPOCHS):
    # Do Training
    total_loss = 0.0
    num_batches = 0
    for batch in train_dist_dataset:
        total_loss += distributed_train_step(batch)
        num_batches += 1
    train_loss = total_loss / num_batches
```



```
def train_step(inputs):  
    images, labels = inputs  
    with tf.GradientTape() as tape:  
        predictions = model(images, training=True)  
        loss = compute_loss(labels, predictions)  
  
    gradients = tape.gradient(loss, model.trainable_variables)  
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))  
  
    train_accuracy.update_state(labels, predictions)  
    return loss
```

```
def train_step(inputs):  
    images, labels = inputs  
    with tf.GradientTape() as tape:  
        predictions = model(images, training=True)  
        loss = compute_loss(labels, predictions)  
  
    gradients = tape.gradient(loss, model.trainable_variables)  
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))  
  
    train_accuracy.update_state(labels, predictions)  
    return loss
```

```
def train_step(inputs):  
    images, labels = inputs  
    with tf.GradientTape() as tape:  
        predictions = model(images, training=True)  
        loss = compute_loss(labels, predictions)  
  
    gradients = tape.gradient(loss, model.trainable_variables)  
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))  
  
    train_accuracy.update_state(labels, predictions)  
    return loss
```



```
def train_step(inputs):  
    images, labels = inputs  
    with tf.GradientTape() as tape:  
        predictions = model(images, training=True)  
        loss = compute_loss(labels, predictions)  
  
        gradients = tape.gradient(loss, model.trainable_variables)  
        optimizer.apply_gradients(zip(gradients, model.trainable_variables))  
  
    train_accuracy.update_state(labels, predictions)  
    return loss
```

```
def train_step(inputs):  
    images, labels = inputs  
    with tf.GradientTape() as tape:  
        predictions = model(images, training=True)  
        loss = compute_loss(labels, predictions)  
  
    gradients = tape.gradient(loss, model.trainable_variables)  
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))  
  
    train_accuracy.update_state(labels, predictions)  
    return loss
```

```
def train_step(inputs):  
    images, labels = inputs  
    with tf.GradientTape() as tape:  
        predictions = model(images, training=True)  
        loss = compute_loss(labels, predictions)  
  
    gradients = tape.gradient(loss, model.trainable_variables)  
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))  
  
    train_accuracy.update_state(labels, predictions)  
    return loss
```


Notebook settings

Hardware accelerator

TPU



To get the most out of Colab Pro, avoid using a TPU unless you need one. [Learn more](#)

Runtime shape

Standard



☐ Omit code cell output when saving this notebook

CANCEL

SAVE

```
# Detect hardware
```

```
try:
```

```
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
```

```
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address)
```

```
    tf.config.experimental_connect_to_cluster(tpu)
```

```
    tf.tpu.experimental.initialize_tpu_system(tpu)
```

```
    strategy = tf.distribute.experimental.TPUStrategy(tpu)
```

```
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
```

```
    print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
except ValueError:
```

```
    print('TPU failed to initialize.')
```

```
# Detect hardware
```

```
try:
```

```
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']  
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address)  
    tf.config.experimental_connect_to_cluster(tpu)  
    tf.tpu.experimental.initialize_tpu_system(tpu)  
    strategy = tf.distribute.experimental.TPUStrategy(tpu)  
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])  
    print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
except ValueError:
```

```
    print('TPU failed to initialize.')
```

```
# Detect hardware
```

```
try:
```

```
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
```

```
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address)
```

```
    tf.config.experimental_connect_to_cluster(tpu)
```

```
    tf.tpu.experimental.initialize_tpu_system(tpu)
```

```
    strategy = tf.distribute.experimental.TPUStrategy(tpu)
```

```
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
```

```
    print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
except ValueError:
```

```
    print('TPU failed to initialize.')
```



```
# Detect hardware
```

```
try:
```

```
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
```

```
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address)
```

```
    tf.config.experimental_connect_to_cluster(tpu)
```

```
    tf.tpu.experimental.initialize_tpu_system(tpu)
```

```
    strategy = tf.distribute.experimental.TPUStrategy(tpu)
```

```
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
```

```
    print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
except ValueError:
```

```
    print('TPU failed to initialize.')
```

```
# Detect hardware
```

```
try:
```

```
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
```

```
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address)
```

```
    tf.config.experimental_connect_to_cluster(tpu)
```

```
    tf.tpu.experimental.initialize_tpu_system(tpu)
```

```
    strategy = tf.distribute.experimental.TPUStrategy(tpu)
```

```
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
```

```
    print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
except ValueError:
```

```
    print('TPU failed to initialize.')
```

```
# Detect hardware
```

```
try:
```

```
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
```

```
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address)
```

```
    tf.config.experimental_connect_to_cluster(tpu)
```

```
    tf.tpu.experimental.initialize_tpu_system(tpu)
```

```
    strategy = tf.distribute.experimental.TPUStrategy(tpu)
```

```
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
```

```
    print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
except ValueError:
```

```
    print('TPU failed to initialize.')
```

```
# Detect hardware
```

```
try:
```

```
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
```

```
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address)
```

```
    tf.config.experimental_connect_to_cluster(tpu)
```

```
    tf.tpu.experimental.initialize_tpu_system(tpu)
```

```
    strategy = tf.distribute.experimental.TPUStrategy(tpu)
```

```
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
```

```
    print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
except ValueError:
```

```
    print('TPU failed to initialize.')
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:TPU_SYSTEM:0,  
TPU_SYSTEM, 0, 0)
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:TPU_SYSTEM:0,  
TPU_SYSTEM, 0, 0)
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:XLA_CPU:0,  
XLA_CPU, 0, 0)
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:XLA_CPU:0,  
XLA_CPU, 0, 0)
```

```
Running on TPU ['10.109.132.10:8470']  
Number of accelerators: 8
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:TPU_SYSTEM:0,  
TPU_SYSTEM, 0, 0)
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:TPU_SYSTEM:0,  
TPU_SYSTEM, 0, 0)
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:XLA_CPU:0,  
XLA_CPU, 0, 0)
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:XLA_CPU:0,  
XLA_CPU, 0, 0)
```

```
Running on TPU ['10.109.132.10:8470']  
Number of accelerators: 8
```

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in `per-replica-losses` structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in `per-replica-losses` structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in `per-replica-losses` structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in per-replica-losses structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in `per-replica-losses` structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in `per-replica-losses` structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in `per-replica-losses` structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in `per-replica-losses` structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

```
@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses, axis=None)

def train_step(inputs):
    images, labels = inputs
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    train_accuracy.update_state(labels, predictions)
    return loss
```

```
@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses, axis=None)
```

```
def train_step(inputs):
    images, labels = inputs
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    train_accuracy.update_state(labels, predictions)
    return loss
```



```
@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses, axis=None)

def train_step(inputs):
    images, labels = inputs
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    train_accuracy.update_state(labels, predictions)
    return loss
```

```
@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses, axis=None)

def train_step(inputs):
    images, labels = inputs
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    train_accuracy.update_state(labels, predictions)
    return loss
```

```
@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses, axis=None)

def train_step(inputs):
    images, labels = inputs
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    train_accuracy.update_state(labels, predictions)
    return loss
```

```
@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses, axis=None)

def train_step(inputs):
    images, labels = inputs
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)
        gradients = tape.gradient(loss, model.trainable_variables)
        optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    train_accuracy.update_state(labels, predictions)
    return loss
```

```
@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses, axis=None)

def train_step(inputs):
    images, labels = inputs
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    train_accuracy.update_state(labels, predictions)
    return loss
```

Training on a single device

```
tf.distribute.OneDeviceStrategy
```



Input data is distributed

```
strategy = tf.distribute.OneDeviceStrategy(device="/gpu:0")
```

Training across many machines

`tf.distribute.experimental.MultiWorkerMirroredStrategy`

- Done across multiple workers, each with multiple GPUs
- Variables are replicated on each device across workers
- Fault tolerance with `tf.keras.callbacks.ModelCheckpoint`
- **Synchronization** done with CollectiveOps



```
multiworker_strategy = tf.distribute.experimental.MultiWorkerMirroredStrategy()
```

Multi-worker training

- Run **workers** in a cluster
- **Tasks** (training/input pipelines)
- **Roles** (chief, worker, ps, evaluator)
- Configuring the cluster (next..)

Cluster specification

```
os.environ["TF_CONFIG"] = json.dumps({  
    "cluster": {  
        "worker": ["host1:port", "host2:port", ...],  
    },  
    "task": {"type": "worker", "index": 0}  
})
```

https://www.tensorflow.org/tutorials/distribute/multi_worker_with_keras

Other strategies

CentralStorageStrategy

Variables - not mirrored, but placed on CPU

Computation - replicated across local GPUs

ParameterServerStrategy

Some machines are designated as workers

... some as parameter servers

Variables - placed on one parameter server (ps)

Computation - replicated across GPUs of all the workers

Other strategies

CentralStorageStrategy

Variables - not mirrored, but placed on CPU

Computation - replicated across local GPUs

ParameterServerStrategy

Some machines are designated as workers

... some as parameter servers

Variables - placed on one parameter server (ps)

Computation - replicated across GPUs of all the workers